

A Simple Approach to Continual Learning By Transferring Skill Parameters

K.R. Zentner, Ryan Julian, Ujjwal Puri,
Yulun Zhang, and Gaurav S. Sukhatme

2021

Bardia Mojra
January 26, 2022
Robotic Vision Lab
University of Texas at Arlington




Introduction

- Problem:
 - existing transfer learning methods are too inefficient for vision-motor tasks (robotic manipulation)
 - Storing past experiences entirely, takes too much space and recall time
- Proposed:
 - Store previous knowledge only in terms of skill policies (parameters)
 - Use general purpose latent state skill representations for decomposed skills
 - Construct more complex skills based on simpler and prior skills
 - Use the propose optimal curriculum learning method for efficient transfer

Meta-World

- › Web-based simulation environment for training DRL models for robotic tasks



button-press



door-open




drawer-close




drawer-open




peg-insert-side



pick-place



push



reach



window-open



window-close

Transfer Learning

- › Defined as learning by transferring knowledge (known skills) from another agent
- › Problems:
 - Traceability problem
 - Reward function estimation and tuning problem
 - Bad transfer

Continual Learning

- › Defined as learning new skills either through real-world experimentation or transferring and re-tuning in simulation
- › Problems:
 - Multi-task learning: could destroy prior skill (pessimal example)
 - Catastrophic forgetting
 - Current methods are too inefficient

Reusable Skill Libraries (DRL-Based Models)

- › Associative Skill Memories
- › Probabilistic Movement Primitive
- › Latent space parameter decomposition

Continual Learning with Skill Libraries and Curricula

- › Learn skills in form of factorized policy model classes
- › Train an online model-based planner for reusing skills with high level action space (domain) for hierarchical RL
- › Use on-policy RL to directly update skills

Setting

- › T: a possibly-unbounded discrete space of tasks.
- › S: a single continuous state space shared among all tasks T.
- › A: a single continuous action space shared among all tasks T.
- › The MTRL problem: (T , S, A), and each task $\tau \in T$ is an infinite-horizon Markov decision process (MDP)

$$\tau = (S, A, M_i, p_{\tau_i}(s, a, s'), r_{\tau_i}(s, a, s')) ,$$

- › M_i : represents the set is manipulation skills the robot is initialized or pre-trained with
- › i: ith the epoch number

Simple Continual Learning with Skill Transfer

Algorithm 1 Proposed Continual Learning Framework

```
1: Input: Initial skill library  $\mathcal{M}_0$ , target task space  $\mathcal{T}$ , RL algorithm  $\mathcal{F} \rightarrow (\pi, \rho)$ , target task rule  
   ChooseTargetTask, base skill rule ChooseBaseSkill  
2:  $i \leftarrow 1$   
3: while not done do  
4:    $\tau \leftarrow \text{ChooseTargetTask}(\mathcal{T}, \mathcal{M}_{i-1})$   
5:   while  $\pi_{\text{target}}$  not solved do  
6:      $\pi_{\text{base}} \leftarrow \text{ChooseBaseSkill}(\mathcal{T}, \mathcal{M}_{i-1})$   
7:      $\pi_{\text{target}}, \cdot \leftarrow \mathcal{F}(\tau, \text{clone}(\pi_{\text{base}}))$   
8:   end while  
9:    $\mathcal{M}_i \leftarrow \{\pi_{\text{target}}\} \cup \mathcal{M}_{i-1}$   
10:   $i \leftarrow i + 1$   
11: end while  
12: Output: Skill library  $\mathcal{M}_i$ 
```

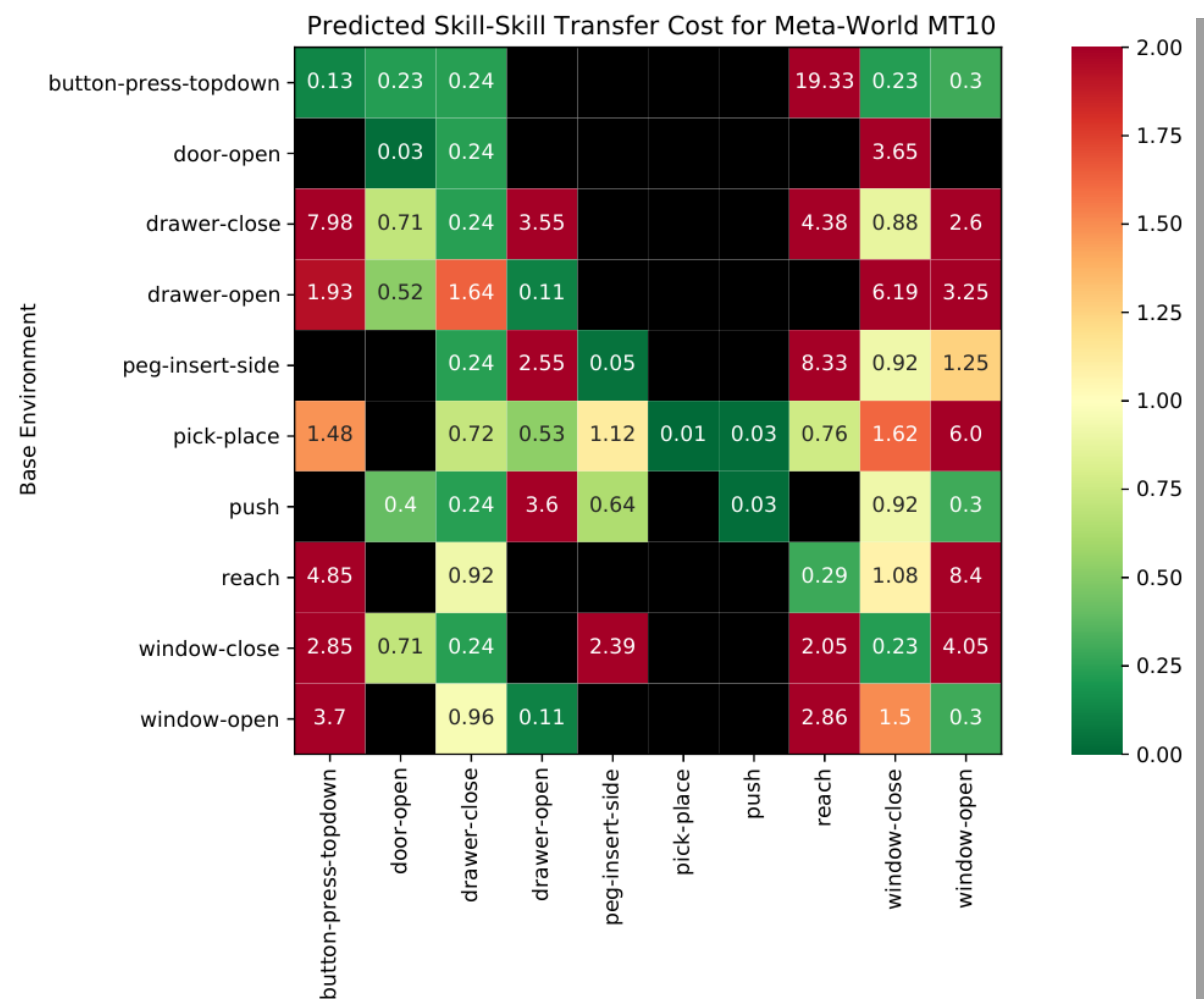
Other Techniques

- › Warm-up procedure for value function transfer
- › Rejecting bad transfers
- › Skill-skill transfer cost
 - Number of samples needed to acquire target skill
- › Predicted skill-skill transfer cost

$$A_{base \rightarrow target} = \frac{C_{base \rightarrow target}}{C_{scratch \rightarrow target}}$$

π

Predicted Skill-Skill Transfer Cost



Curriculum Selection Algorithm

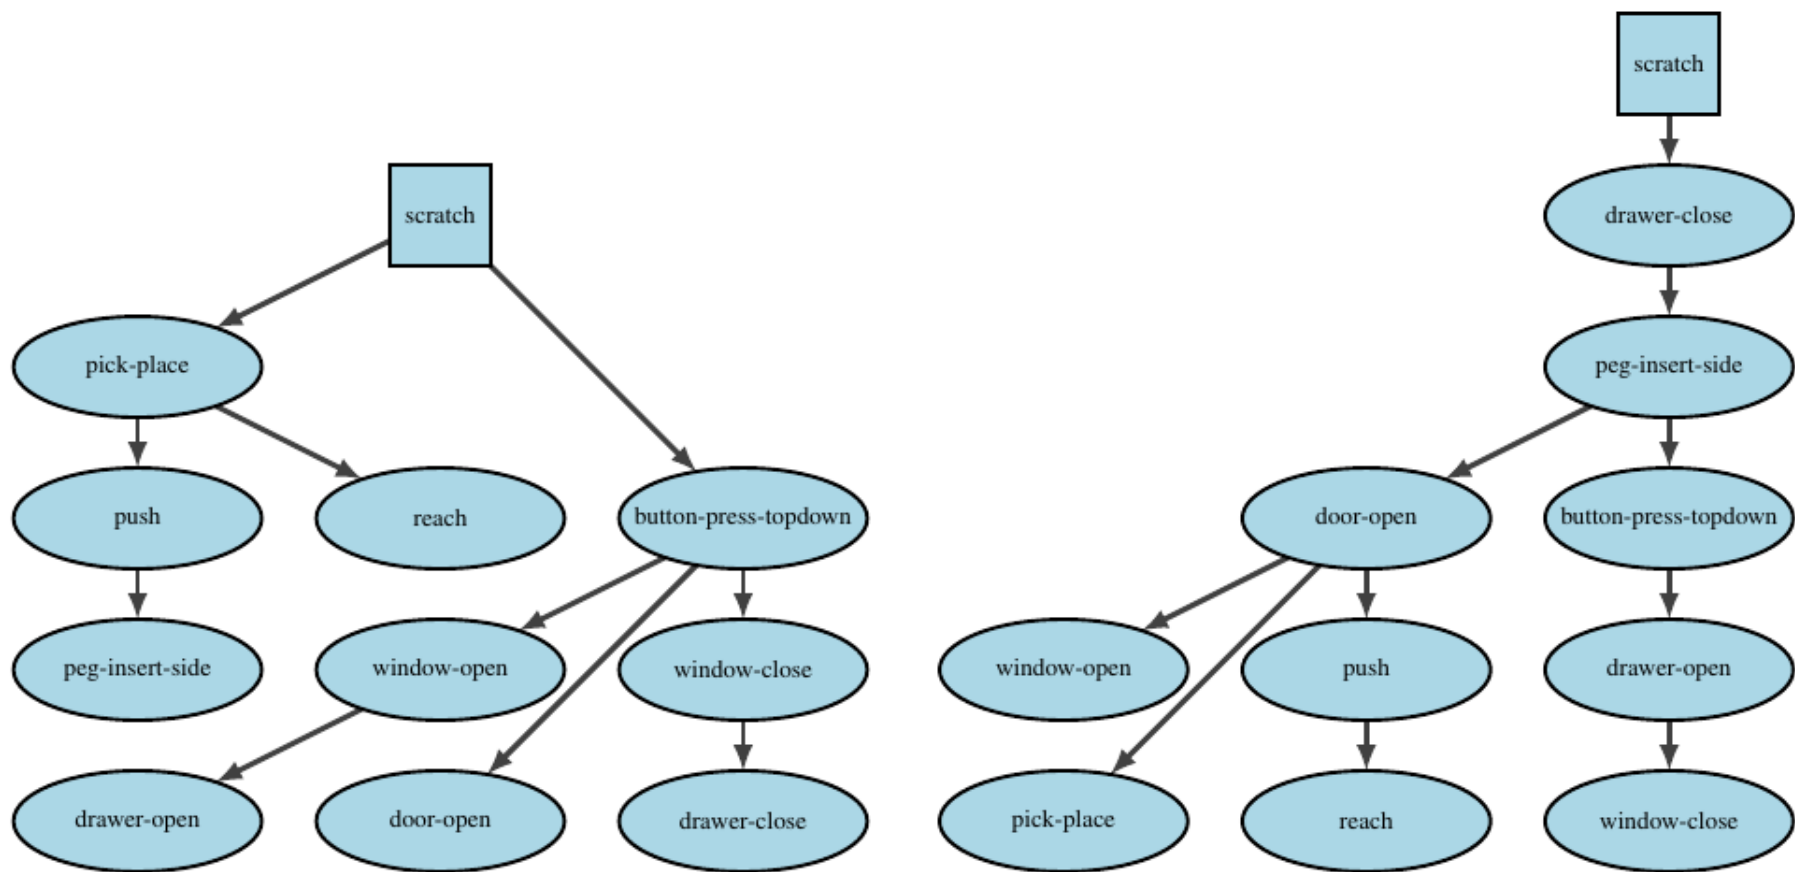
Algorithm 2 DMST-Based Curriculum Transfer

```

1: Input: Initial skill library  $\mathcal{M}_0$ , target task space  $\mathcal{T}$ , RL algorithm  $\mathcal{F} \rightarrow (\pi, C)$ 
2:  $V \leftarrow \mathcal{T} \cup \text{scratch}$ 
3:  $E \leftarrow \{\}$ 
4: for  $\tau_{base} \in \mathcal{T}$  do
5:    $E \leftarrow (\text{scratch}, \tau_{base}, -1.0)$ 
6:    $\pi_{base}, C_{\text{scratch} \rightarrow \text{target}} \leftarrow \mathcal{F}(\tau_{base}, \pi_{\text{random}})$ 
7:   for  $\tau_{target} \in \mathcal{T}$  do
8:      $\cdot, C_{\text{base} \rightarrow \text{target}} = \mathcal{F}(\tau_{target}, \pi_{base})$ 
9:      $E \leftarrow E(\tau_{base}, \tau_{target}, C_{\text{base} \rightarrow \text{target}}) \cup E$ 
10:  end for
11: end for
12:  $T_{\text{optimal}} \leftarrow \text{kruskal}((V, E))$ 
13:  $i \leftarrow 1$ 
14:  $\pi_{base} \leftarrow \pi_{\text{random}}$ 
15: for  $\tau_{target} \in \text{traverse}(T_{\text{optimal}})$  do
16:   while  $\tau_{target}$  not solved do
17:      $\pi_{target}, \cdot \leftarrow \mathcal{F}(\tau_{target}, \text{clone}(\pi_{base}))$ 
18:     if  $\tau_{target}$  not solved then
19:        $E \leftarrow E \setminus (\tau_{base}, \tau_{target})$ 
20:        $T_{\text{optimal}} \leftarrow \text{kruskal}((V, E))$ 
21:     end if
22:   end while
23:    $\mathcal{M}_i \leftarrow \{\pi_{target}\} \cup \mathcal{M}_{i-1}$ 
24:    $i \leftarrow i + 1$ 
25: end for
26: Output: Skill library  $\mathcal{M}$ 


```

Optimal and Pessimal Curricula



π

Experiments



› Thank you!